

Carrefour Kenya Analysis

Naomi Chebet

9/14/2020

1. Business Understanding

1.1 Define the Question

As a Data analyst at Carrefour Kenya, you are tasked with conducting analysis of sales data to inform the marketing department on the most relevant marketing strategies that will result in the highest no. of sales (total price including tax). In this project you are required to implement PCA and Feature Selection methods.

1.2 Understanding the Context

Carrefour is a French Multinational Corporation that specializes in retail. In 2016 Carrefour opened it's stores in Kenya to what is now called Carrefour Kenya. Carrefour Kenya competes with other big retail shops such as Quickmart and Naivas. As a data scientist, you've been provided with sales data to conduct analysis, the results from the analysis will inform the marketing team on the best strategies that will lead to high number of sales.

The dataset provided has 1000 entries and 16 columns, some of the columns in the dataset are sales, Branch, Customer.type, Gender, Product.line, Quantity, and Date.

1.3 Metrics of Success

1. Successful implementation of PCA and obtaining insights from the analysis
2. Successful implementation of at least two feature selection methods
3. Providing recommendation based on EDA and solution implementation analysis

1.4 Experimental Design

The flow of our project includes:

1. Business Understanding
2. Loading and Checking the Data
3. Tidying the Data
4. EDA with Univariate, and Bivariate Analysis
5. Implementing the solution with PCA and Feature Selection Methods

2. Importing libraries

```
library(tidyverse)
```

```
## -- Attaching packages -----
```

```

## v ggplot2 3.3.2      v purrr  0.3.4
## v tibble  3.0.3      v dplyr  1.0.2
## v tidyr   1.1.2      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.5.0

## -- Conflicts -----
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(ggbiplot)

## Loading required package: plyr
## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
## -----
##
## Attaching package: 'plyr'
##
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
##
## The following object is masked from 'package:purrr':
##
##   compact
##
## Loading required package: scales
##
## Attaching package: 'scales'
##
## The following object is masked from 'package:purrr':
##
##   discard
##
## The following object is masked from 'package:readr':
##
##   col_factor
##
## Loading required package: grid

library(ggplot2)
library(lubridate)

##
## Attaching package: 'lubridate'
##
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

library(devtools)

## Loading required package: usethis

```

```

library(tidyr)
library(corrplot)

## corrplot 0.84 loaded

library(ggcorrplot)
library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
## The following object is masked from 'package:purrr':
##
##     lift
library(clustvarsel)

## Loading required package: mclust
## Package 'mclust' version 5.4.6
## Type 'citation("mclust")' for citing this R package in publications.
##
## Attaching package: 'mclust'
## The following object is masked from 'package:purrr':
##
##     map
## Package 'clustvarsel' version 2.3.3
## Type 'citation("clustvarsel")' for citing this R package in publications.
library(mclust)
library(wskm)

## Loading required package: latticeExtra
##
## Attaching package: 'latticeExtra'
## The following object is masked from 'package:ggplot2':
##
##     layer
## Loading required package: fpc
library(cluster)
# library(FSelector)

```

3. Loading and Checking the Data

```

sales <- read.csv("~/Moringa School/R Programming/R datasets/Supermarket_Dataset_1 - Sales Data.csv")
glimpse(sales)

## Rows: 1,000
## Columns: 16
## $ Invoice.ID      <chr> "750-67-8428", "226-31-3081", "631-41-3108"...

```

```
## $ Branch      <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ Customer.type <chr> "Member", "Normal", "Normal", "Member", "No...
## $ Gender      <chr> "Female", "Female", "Male", "Male", "Male",...
## $ Product.line <chr> "Health and beauty", "Electronic accessorie...
## $ Unit.price   <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ Quantity     <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ Tax          <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ Date         <chr> "1/5/2019", "3/8/2019", "3/3/2019", "1/27/2...
## $ Time         <chr> "13:08", "10:29", "13:23", "20:33", "10:37"...
## $ Payment      <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
## $ cogs         <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.income <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ Rating       <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ Total        <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
```

```
class(sales)
```

```
## [1] "data.frame"
```

```
# previewing the top of the data
```

```
head(sales)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 1 750-67-8428      A      Member Female      Health and beauty      74.69
## 2 226-31-3081      C      Normal Female Electronic accessories      15.28
## 3 631-41-3108      A      Normal  Male      Home and lifestyle      46.33
## 4 123-19-1176      A      Member  Male      Health and beauty      58.22
## 5 373-73-7910      A      Normal  Male      Sports and travel      86.31
## 6 699-14-3026      C      Normal  Male      Electronic accessories      85.39
##      Quantity      Tax      Date Time      Payment      cogs gross.margin.percentage
## 1          7 26.1415 1/5/2019 13:08      Ewallet 522.83          4.761905
## 2          5  3.8200 3/8/2019 10:29      Cash 76.40          4.761905
## 3          7 16.2155 3/3/2019 13:23 Credit card 324.31          4.761905
## 4          8 23.2880 1/27/2019 20:33      Ewallet 465.76          4.761905
## 5          7 30.2085 2/8/2019 10:37      Ewallet 604.17          4.761905
## 6          7 29.8865 3/25/2019 18:30      Ewallet 597.73          4.761905
##      gross.income Rating      Total
## 1      26.1415      9.1 548.9715
## 2       3.8200      9.6  80.2200
## 3      16.2155      7.4 340.5255
## 4      23.2880      8.4 489.0480
## 5      30.2085      5.3 634.3785
## 6      29.8865      4.1 627.6165
```

```
# previewing the bottom
```

```
tail(sales)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 995 652-49-6720      C      Member Female Electronic accessories      60.95
## 996 233-67-5758      C      Normal  Male      Health and beauty      40.35
## 997 303-96-2227      B      Normal Female      Home and lifestyle      97.38
## 998 727-02-1313      A      Member  Male      Food and beverages      31.84
## 999 347-56-2442      A      Normal  Male      Home and lifestyle      65.82
## 1000 849-09-3807      A      Member Female      Fashion accessories      88.34
##      Quantity      Tax      Date Time Payment      cogs gross.margin.percentage
## 995          1  3.0475 2/18/2019 11:40 Ewallet 60.95          4.761905
```

```
## 996      1  2.0175 1/29/2019 13:46 Ewallet  40.35      4.761905
## 997     10 48.6900 3/2/2019 17:16 Ewallet 973.80      4.761905
## 998      1  1.5920 2/9/2019 13:22   Cash  31.84      4.761905
## 999      1  3.2910 2/22/2019 15:33   Cash  65.82      4.761905
## 1000     7 30.9190 2/18/2019 13:28   Cash 618.38      4.761905
##      gross.income Rating      Total
## 995      3.0475      5.9    63.9975
## 996      2.0175      6.2    42.3675
## 997     48.6900      4.4  1022.4900
## 998      1.5920      7.7    33.4320
## 999      3.2910      4.1    69.1110
## 1000     30.9190      6.6   649.2990
```

```
# checking column names
```

```
colnames(sales)
```

```
## [1] "Invoice.ID"      "Branch"
## [3] "Customer.type"   "Gender"
## [5] "Product.line"    "Unit.price"
## [7] "Quantity"        "Tax"
## [9] "Date"            "Time"
## [11] "Payment"         "cogs"
## [13] "gross.margin.percentage" "gross.income"
## [15] "Rating"          "Total"
```

```
# checking for the total number of unique values in each column
```

```
sapply(sales, function(x)length(unique(x)))
```

```
##      Invoice.ID      Branch      Customer.type
##      1000          3          2
##      Gender      Product.line      Unit.price
##      2          6          943
##      Quantity      Tax      Date
##      10          990      89
##      Time      Payment      cogs
##      506          3      990
## gross.margin.percentage      gross.income      Rating
##      1          990      61
##      Total
##      990
```

```
# displaying the unique values in columns of interest
```

```
unique.df <- select(sales, Branch, Customer.type, Gender, Product.line, Quantity, Date,
                    Payment, gross.margin.percentage, Rating)
```

```
un <- lapply(unique.df, unique)
```

```
un
```

```
## $Branch
```

```
## [1] "A" "C" "B"
```

```
##
```

```
## $Customer.type
```

```
## [1] "Member" "Normal"
```

```
##
```

```
## $Gender
```

```
## [1] "Female" "Male"
```

```
##
```

```
## $Product.line
```

```
## [1] "Health and beauty"      "Electronic accessories" "Home and lifestyle"
## [4] "Sports and travel"      "Food and beverages"    "Fashion accessories"
##
## $Quantity
## [1] 7 5 8 6 10 2 3 4 1 9
##
## $Date
## [1] "1/5/2019" "3/8/2019" "3/3/2019" "1/27/2019" "2/8/2019" "3/25/2019"
## [7] "2/25/2019" "2/24/2019" "1/10/2019" "2/20/2019" "2/6/2019" "3/9/2019"
## [13] "2/12/2019" "2/7/2019" "3/29/2019" "1/15/2019" "3/11/2019" "1/1/2019"
## [19] "1/21/2019" "3/5/2019" "3/15/2019" "2/17/2019" "3/2/2019" "3/22/2019"
## [25] "3/10/2019" "1/25/2019" "1/28/2019" "1/7/2019" "3/23/2019" "1/17/2019"
## [31] "2/2/2019" "3/4/2019" "3/16/2019" "2/27/2019" "2/10/2019" "3/19/2019"
## [37] "2/3/2019" "3/7/2019" "2/28/2019" "3/27/2019" "1/20/2019" "3/12/2019"
## [43] "2/15/2019" "3/6/2019" "2/14/2019" "3/13/2019" "1/24/2019" "1/6/2019"
## [49] "2/11/2019" "1/22/2019" "1/13/2019" "1/9/2019" "1/12/2019" "1/26/2019"
## [55] "1/23/2019" "2/23/2019" "1/2/2019" "2/9/2019" "3/26/2019" "3/1/2019"
## [61] "2/1/2019" "3/28/2019" "3/24/2019" "2/5/2019" "1/19/2019" "1/16/2019"
## [67] "1/8/2019" "2/18/2019" "1/18/2019" "2/16/2019" "2/22/2019" "1/29/2019"
## [73] "1/4/2019" "3/30/2019" "1/30/2019" "1/3/2019" "3/21/2019" "2/13/2019"
## [79] "1/14/2019" "3/18/2019" "3/20/2019" "2/21/2019" "1/31/2019" "1/11/2019"
## [85] "2/26/2019" "3/17/2019" "3/14/2019" "2/4/2019" "2/19/2019"
##
## $Payment
## [1] "Ewallet" "Cash" "Credit card"
##
## $gross.margin.percentage
## [1] 4.761905
##
## $Rating
## [1] 9.1 9.6 7.4 8.4 5.3 4.1 5.8 8.0 7.2 5.9 4.5 6.8 7.1 8.2 5.7
## [16] 4.6 6.9 8.6 4.4 4.8 5.1 9.9 6.0 8.5 6.7 7.7 7.5 7.0 4.7 7.6
## [31] 7.9 6.3 5.6 9.5 8.1 6.5 6.1 6.6 5.4 9.3 10.0 6.4 4.3 4.0 8.7
## [46] 9.4 5.5 8.3 7.3 4.9 4.2 9.2 7.8 5.2 9.0 8.8 6.2 9.8 9.7 5.0
## [61] 8.9
```

There are three branches: A, B, C. Customer type, we have Member and Normal. We have 6 product lines. Quantity, this is the number of items that a customer bought. We have different options from 1 - 10. There are three payment options: Ewallet, Cash or Credit Card The gross margin percent is constant at 4.761905

```
# converting column names to lower case for uniformity purposes
colnames(sales) <- tolower(colnames(sales))
glimpse(sales)
```

```
## Rows: 1,000
## Columns: 16
## $ invoice.id      <chr> "750-67-8428", "226-31-3081", "631-41-3108"...
## $ branch          <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ customer.type    <chr> "Member", "Normal", "Normal", "Member", "No...
## $ gender           <chr> "Female", "Female", "Male", "Male", "Male",...
## $ product.line     <chr> "Health and beauty", "Electronic accessorie...
## $ unit.price       <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity         <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax              <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ date             <chr> "1/5/2019", "3/8/2019", "3/3/2019", "1/27/2..."
```

```
## $ time                <chr> "13:08", "10:29", "13:23", "20:33", "10:37"...
## $ payment             <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
## $ cogs                <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.income        <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating              <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total               <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
```

```
# splitting the Date column to date, month and year
```

```
sales <- sales %>%
  mutate(date = mdy(date)) %>%
  mutate_at(vars(date), funs(month, day, year)) %>%
  glimpse() # to confirm that the data column was successfully split
```

```
## Warning: `funs()` is deprecated as of dplyr 0.8.0.
```

```
## Please use a list of either functions or lambdas:
```

```
##
```

```
##   # Simple named list:
```

```
##   list(mean = mean, median = median)
```

```
##
```

```
##   # Auto named with `tibble::lst()`:
```

```
##   tibble::lst(mean, median)
```

```
##
```

```
##   # Using lambdas
```

```
##   list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
```

```
## This warning is displayed once every 8 hours.
```

```
## Call `lifecycle::last_warnings()` to see where this warning was generated.
```

```
## Rows: 1,000
```

```
## Columns: 19
```

```
## $ invoice.id          <chr> "750-67-8428", "226-31-3081", "631-41-3108"...
## $ branch              <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ customer.type       <chr> "Member", "Normal", "Normal", "Member", "No...
## $ gender              <chr> "Female", "Female", "Male", "Male", "Male",...
## $ product.line        <chr> "Health and beauty", "Electronic accessorie...
## $ unit.price          <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity            <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax                 <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ date                <date> 2019-01-05, 2019-03-08, 2019-03-03, 2019-0...
## $ time                <chr> "13:08", "10:29", "13:23", "20:33", "10:37"...
## $ payment             <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
## $ cogs                <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.income        <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating              <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total               <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
## $ month               <dbl> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day                 <int> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ year                <dbl> 2019, 2019, 2019, 2019, 2019, 2019, 2019, 2...
```

```
# finding the unique values in month, year and day
```

```
unique.date <- select(sales, month, day, year)
```

```
uni.date <- sapply(unique.date, unique)
```

```
uni.date
```

```
## $month
```

```
## [1] 1 3 2
##
## $day
## [1] 5 8 3 27 25 24 10 20 6 9 12 7 29 15 11 1 21 17 2 22 28 23 4 16 19
## [26] 14 13 26 18 30 31
##
## $year
## [1] 2019
```

There are three months in the dataset: January, February and March. The data is from 2019, and all the dates of a month are listed. There are no strange values in the month, day and year columns.

```
# converting the time column into time format
```

```
sales <- sales %>% mutate(time = hm(time))
glimpse(sales)
```

```
## Rows: 1,000
## Columns: 19
## $ invoice.id      <chr> "750-67-8428", "226-31-3081", "631-41-3108"...
## $ branch          <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ customer.type   <chr> "Member", "Normal", "Normal", "Member", "No...
## $ gender           <chr> "Female", "Female", "Male", "Male", "Male",...
## $ product.line     <chr> "Health and beauty", "Electronic accessorie...
## $ unit.price       <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity         <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax              <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ date             <date> 2019-01-05, 2019-03-08, 2019-03-03, 2019-0...
## $ time             <Period> 13H 8M 0S, 10H 29M 0S, 13H 23M 0S, 20H 3...
## $ payment          <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
## $ cogs             <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.income     <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating           <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total            <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
## $ month            <dbl> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day              <int> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ year             <dbl> 2019, 2019, 2019, 2019, 2019, 2019, 2019, 2...
```

```
# splitting time to hour and minute into a new column
```

```
sales <- sales %>% mutate(hour = hour(time), minute = minute(time))
glimpse(sales)
```

```
## Rows: 1,000
## Columns: 21
## $ invoice.id      <chr> "750-67-8428", "226-31-3081", "631-41-3108"...
## $ branch          <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ customer.type   <chr> "Member", "Normal", "Normal", "Member", "No...
## $ gender           <chr> "Female", "Female", "Male", "Male", "Male",...
## $ product.line     <chr> "Health and beauty", "Electronic accessorie...
## $ unit.price       <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity         <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax              <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ date             <date> 2019-01-05, 2019-03-08, 2019-03-03, 2019-0...
## $ time             <Period> 13H 8M 0S, 10H 29M 0S, 13H 23M 0S, 20H 3...
## $ payment          <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
```



```
## $ cogs <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.income <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
## $ month <dbl> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day <int> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ year <dbl> 2019, 2019, 2019, 2019, 2019, 2019, 2019, 2...
## $ hour <dbl> 13, 10, 13, 20, 10, 18, 14, 11, 17, 13, 18,...
## $ minute <dbl> 8, 29, 23, 33, 37, 30, 36, 38, 15, 27, 7, 3...
```

```
# converting day and month, hour and minute to factor
```

```
sales$month <- as.factor(sales$month)
sales$day <- as.factor(sales$day)
sales$hour <- as.factor(sales$hour)
sales$minute <- as.factor(sales$minute)
```

4. Cleaning the data

4.1 Fixing column names and data types

```
# changing some column names so that they make more sense
```

```
names(sales)[12] <- "cost.of.goods.sold"
names(sales)[16] <- "total.sales.plus.tax"
names(sales)[14] <- "gross.profit"
glimpse(sales)
```

```
## Rows: 1,000
## Columns: 21
## $ invoice.id <chr> "750-67-8428", "226-31-3081", "631-41-3108"...
## $ branch <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A"...
## $ customer.type <chr> "Member", "Normal", "Normal", "Member", "No...
## $ gender <chr> "Female", "Female", "Male", "Male", "Male",...
## $ product.line <chr> "Health and beauty", "Electronic accessorie...
## $ unit.price <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ date <date> 2019-01-05, 2019-03-08, 2019-03-03, 2019-0...
## $ time <Period> 13H 8M 0S, 10H 29M 0S, 13H 23M 0S, 20H 3...
## $ payment <chr> "Ewallet", "Cash", "Credit card", "Ewallet"...
## $ cost.of.goods.sold <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.profit <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total.sales.plus.tax <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
## $ month <fct> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day <fct> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ year <dbl> 2019, 2019, 2019, 2019, 2019, 2019, 2019, 2...
## $ hour <fct> 13, 10, 13, 20, 10, 18, 14, 11, 17, 13, 18,...
## $ minute <fct> 8, 29, 23, 33, 37, 30, 36, 38, 15, 27, 7, 3...
```

```
# converting quantity to character datatype
```

```
sales$quantity <- as.character(sales$quantity)
```

4.2 Missing values

```
# checking for missing values
colSums(is.na(sales))
```

```
##          invoice.id          branch          customer.type
##              0              0              0
##          gender      product.line      unit.price
##              0              0              0
##          quantity      tax          date
##              0              0              0
##          time      payment      cost.of.goods.sold
##              0              0              0
## gross.margin.percentage      gross.profit      rating
##              0              0              0
## total.sales.plus.tax      month      day
##              0              0              0
##          year      hour      minute
##              0              0              0
```

The data has no missing values

4.3 Duplicated

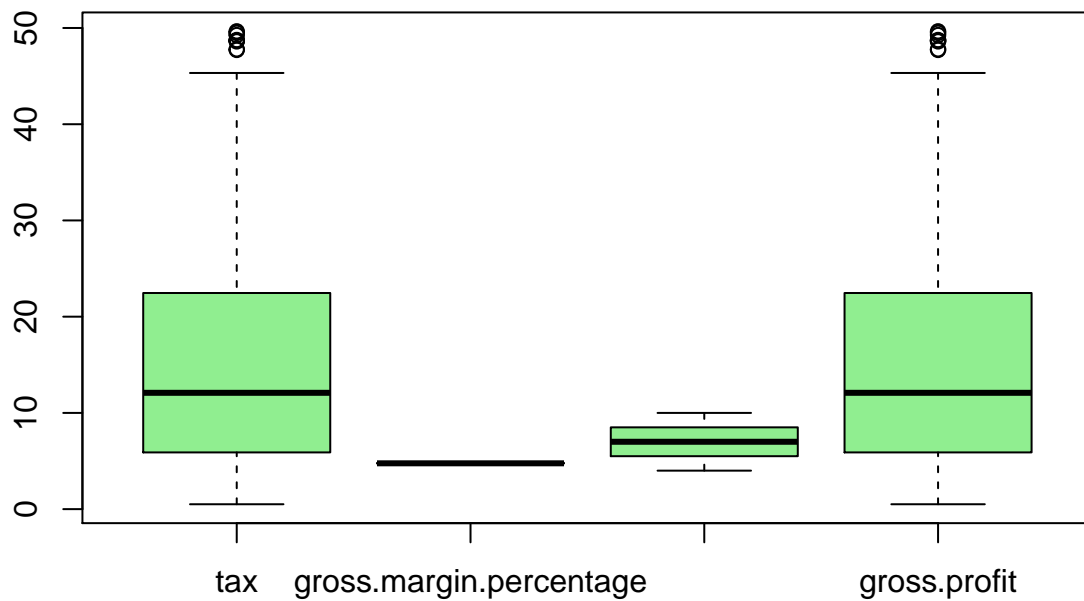
```
dup <- sales[duplicated(sales),]
dup
```

```
## [1] invoice.id          branch          customer.type
## [4] gender      product.line      unit.price
## [7] quantity      tax          date
## [10] time      payment      cost.of.goods.sold
## [13] gross.margin.percentage gross.profit      rating
## [16] total.sales.plus.tax      month      day
## [19] year      hour      minute
## <0 rows> (or 0-length row.names)
```

There are no duplicated records in our dataset

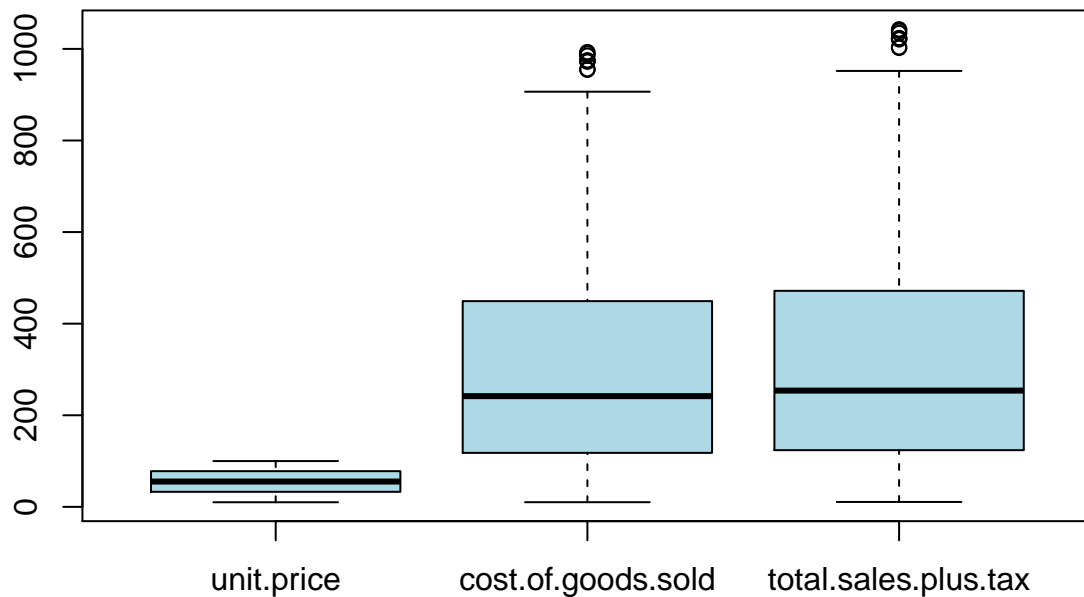
4.4 Outliers

```
boxplot(select(sales, tax, gross.margin.percentage, rating, gross.profit), col = "light green")
```



The data has a few outliers in the tax and gross income columns. We won't be deleting these outliers because we establish that they are true observations and it's possible for some products to be taxed very highly.

```
# more boxplot outliers  
boxplot(select(sales, unit.price, cost.of.goods.sold, total.sales.plus.tax), col = "light blue")
```



There are some outliers in cost of goods sold and total.sales.plus tax columns. This observation is in line with the tax outlier seen earlier. We won't be deleting these outliers because we establish that they are true observations.

4.5 Dropping unnecessary columns

```
# deleting unnecessary column
sales$date <- NULL
sales$year <- NULL
sales$time <- NULL
sales$invoice.id <- NULL
```

5. EDA

5.1 Univariate Analysis

```
# we will start with summary stats for numerical columns
num_col <- select(sales, unit.price, tax, cost.of.goods.sold, gross.profit, rating, total.sales.plus.tax)
summary(num_col)
```

##	unit.price	tax	cost.of.goods.sold	gross.profit
##	Min. :10.08	Min. : 0.5085	Min. : 10.17	Min. : 0.5085
##	1st Qu.:32.88	1st Qu.: 5.9249	1st Qu.:118.50	1st Qu.: 5.9249
##	Median :55.23	Median :12.0880	Median :241.76	Median :12.0880
##	Mean :55.67	Mean :15.3794	Mean :307.59	Mean :15.3794
##	3rd Qu.:77.94	3rd Qu.:22.4453	3rd Qu.:448.90	3rd Qu.:22.4453

```
## Max. :99.96 Max. :49.6500 Max. :993.00 Max. :49.6500
## rating total.sales.plus.tax
## Min. : 4.000 Min. : 10.68
## 1st Qu.: 5.500 1st Qu.: 124.42
## Median : 7.000 Median : 253.85
## Mean : 6.973 Mean : 322.97
## 3rd Qu.: 8.500 3rd Qu.: 471.35
## Max. :10.000 Max. :1042.65
```

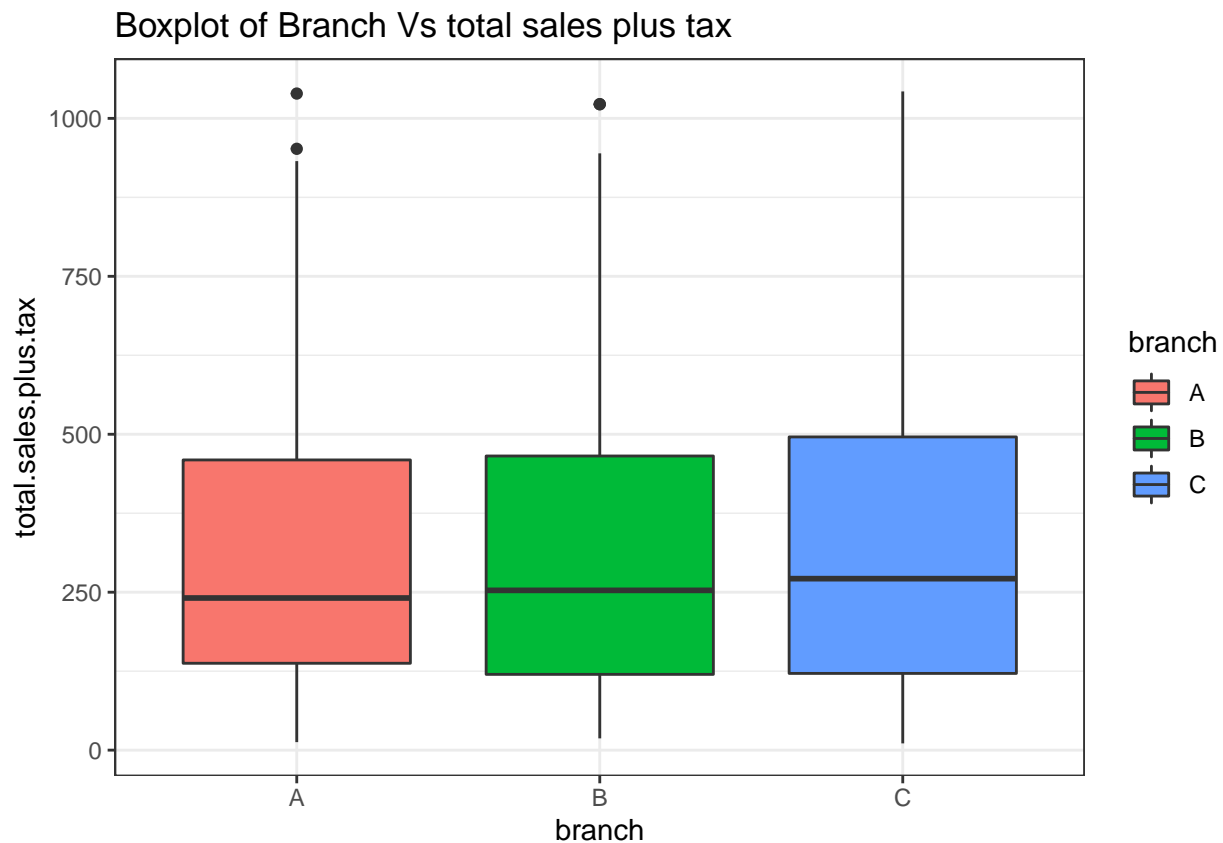
Unit price mean = 55.67 The average quantity of goods sold = 5.51 The mean tax = 15.37 The average cost of goods sold = 307.59 Gross margin is constant at 4.762 Gross profit = 15.37 The average rating is 6.97 The average total sales plus tax = 322.97

5.2 Bivariate Analysis

5.2.1 Boxplots for comparing the target variable with categorical variables

Our target variable is total.sales.plus.tax. Let's do some analysis to see how that compares with oth

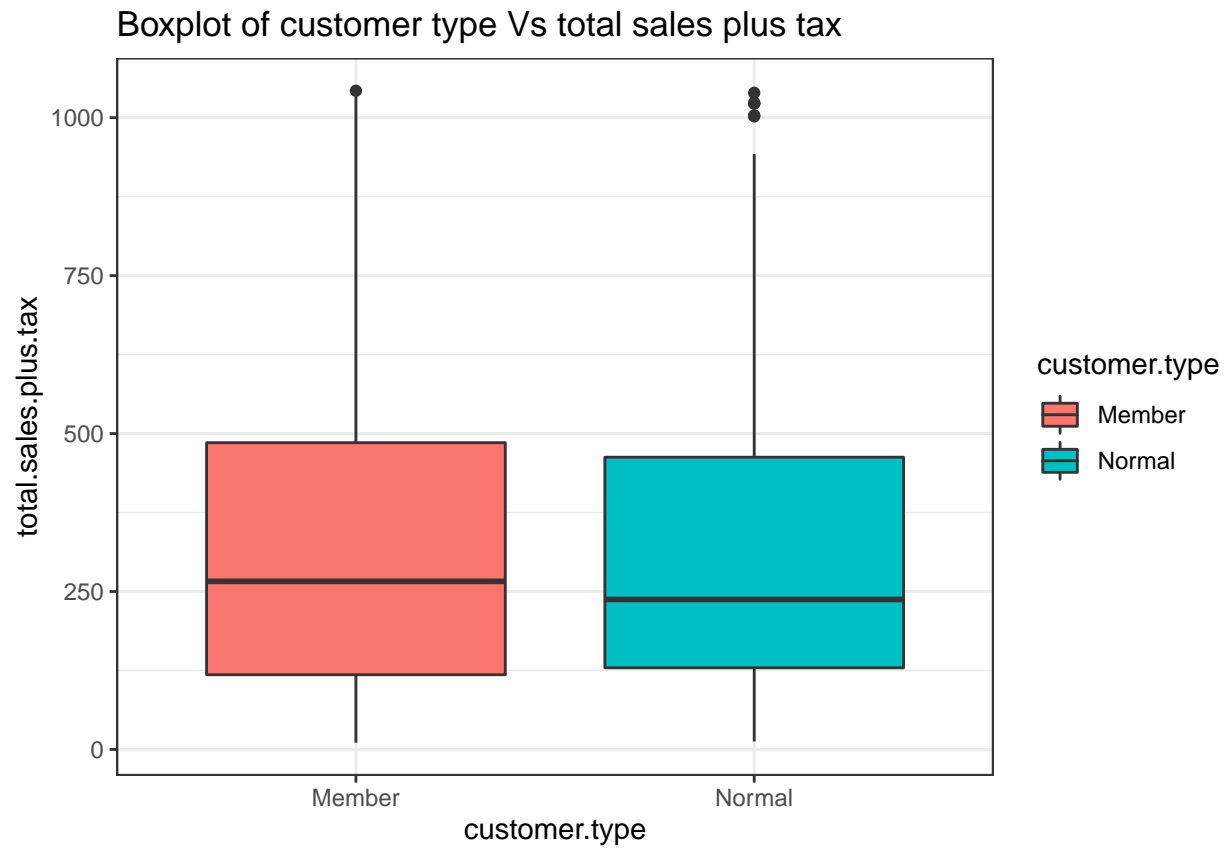
```
ggplot(sales, aes(x=branch, y=total.sales.plus.tax, fill = branch)) +
  theme_bw() +
  geom_boxplot() +
  labs(title = "Boxplot of Branch Vs total sales plus tax")
```



We have a few outliers in branch A and B but none in C. Branch C is bringing in slightly more money.

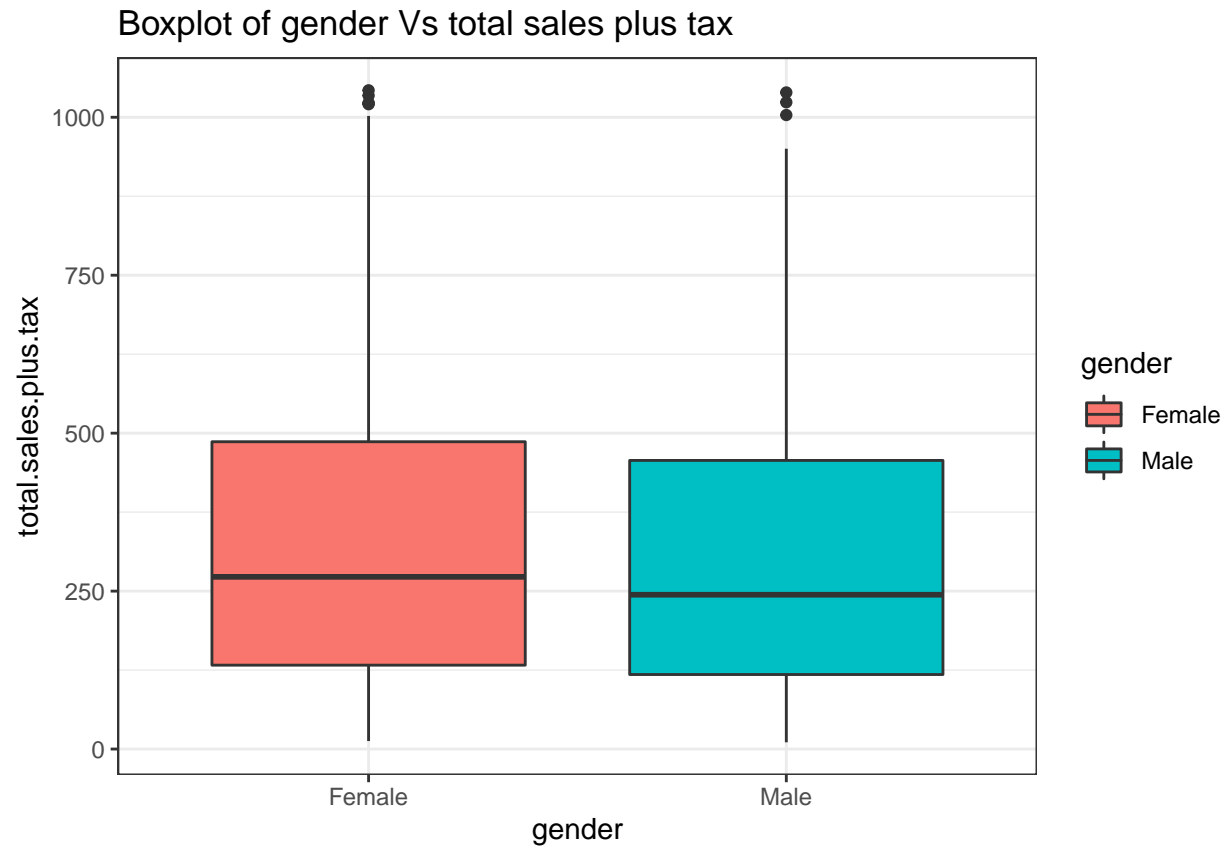
```
ggplot(sales, aes(x=customer.type, y=total.sales.plus.tax, fill = customer.type)) +
  theme_bw() +
```

```
geom_boxplot() +  
labs(title = "Boxplot of customer type Vs total sales plus tax")
```



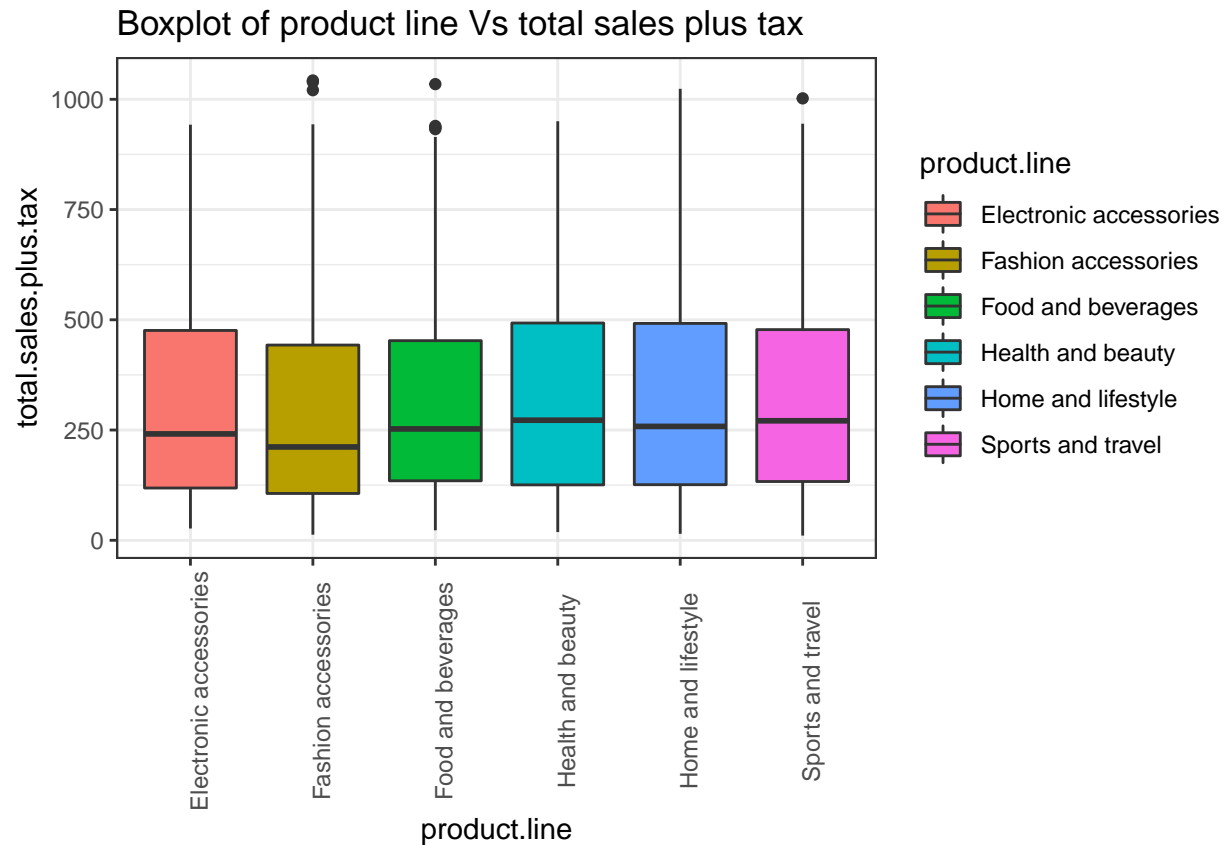
The members customer type on average bring in more revenue through sales.

```
ggplot(sales, aes(x=gender, y=total.sales.plus.tax, fill = gender)) +  
  theme_bw() +  
  geom_boxplot() +  
  labs(title = "Boxplot of gender Vs total sales plus tax")
```



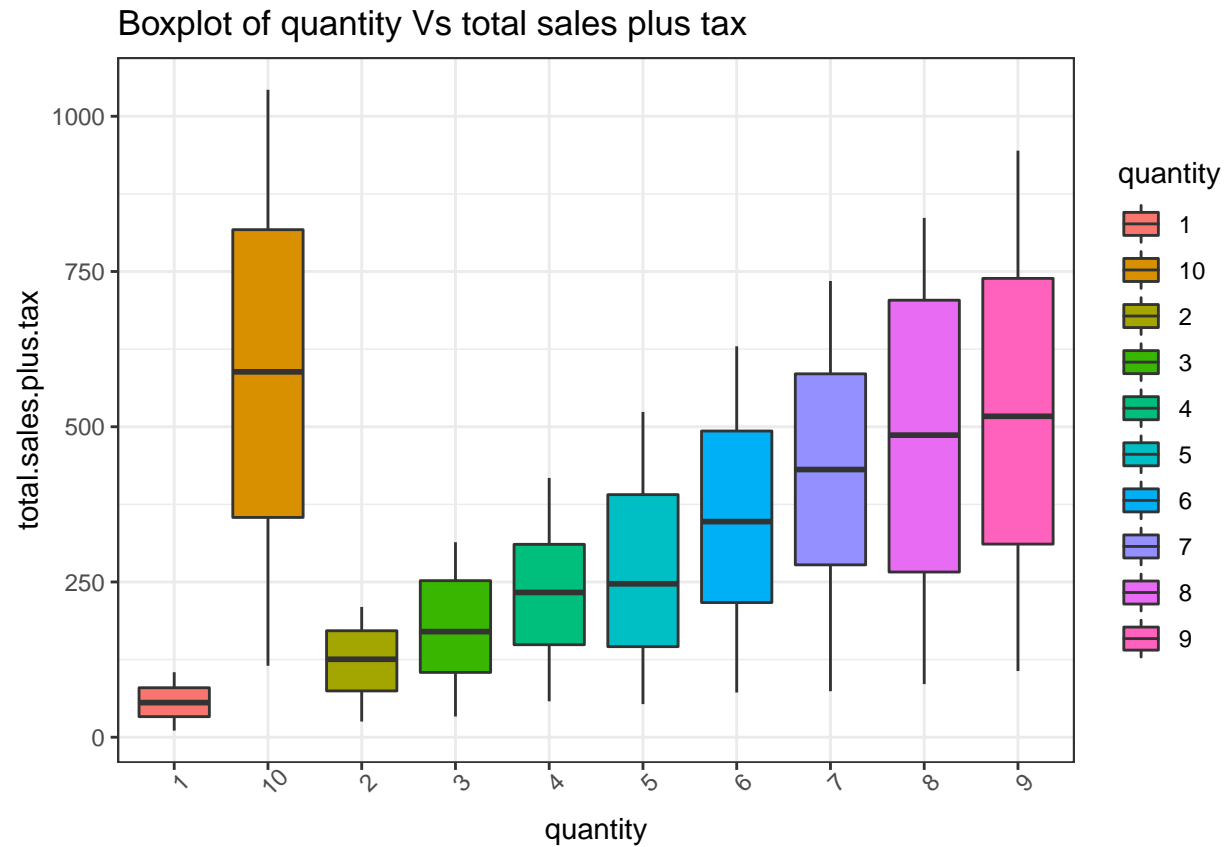
Total sales from women is a bit more compared to that of men.

```
ggplot(sales, aes(x=product.line, y=total.sales.plus.tax, fill = product.line)) +  
  theme_bw() +  
  geom_boxplot() +  
  labs(title = "Boxplot of product line Vs total sales plus tax") +  
  theme(axis.text.x = element_text(angle = 90))
```



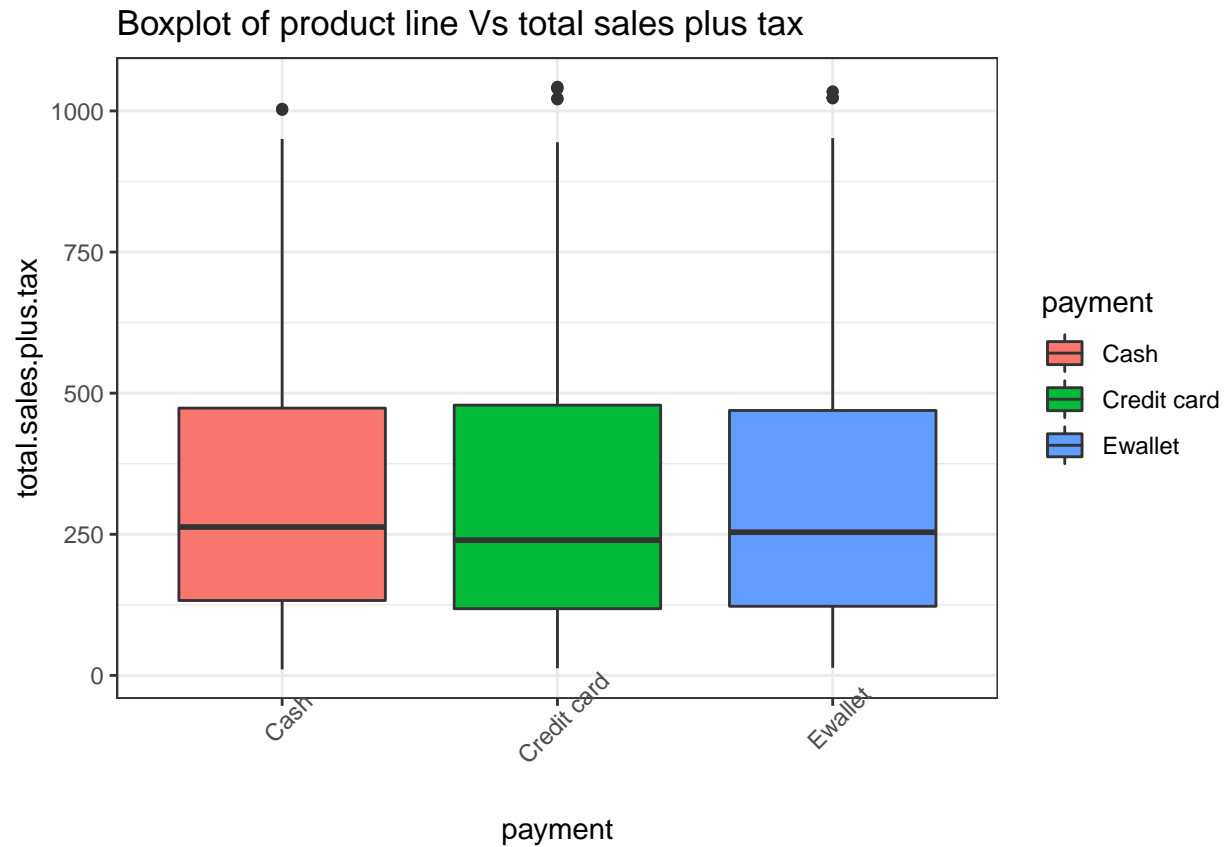
Based on the boxplot, on average, the product line bringing in the most income is health and beauty, and home and lifestyle.

```
ggplot(sales, aes(x=quantity, y=total.sales.plus.tax, fill = quantity)) +
  theme_bw() +
  geom_boxplot() +
  labs(title = "Boxplot of quantity Vs total sales plus tax") +
  theme(axis.text.x = element_text(angle = 45))
```

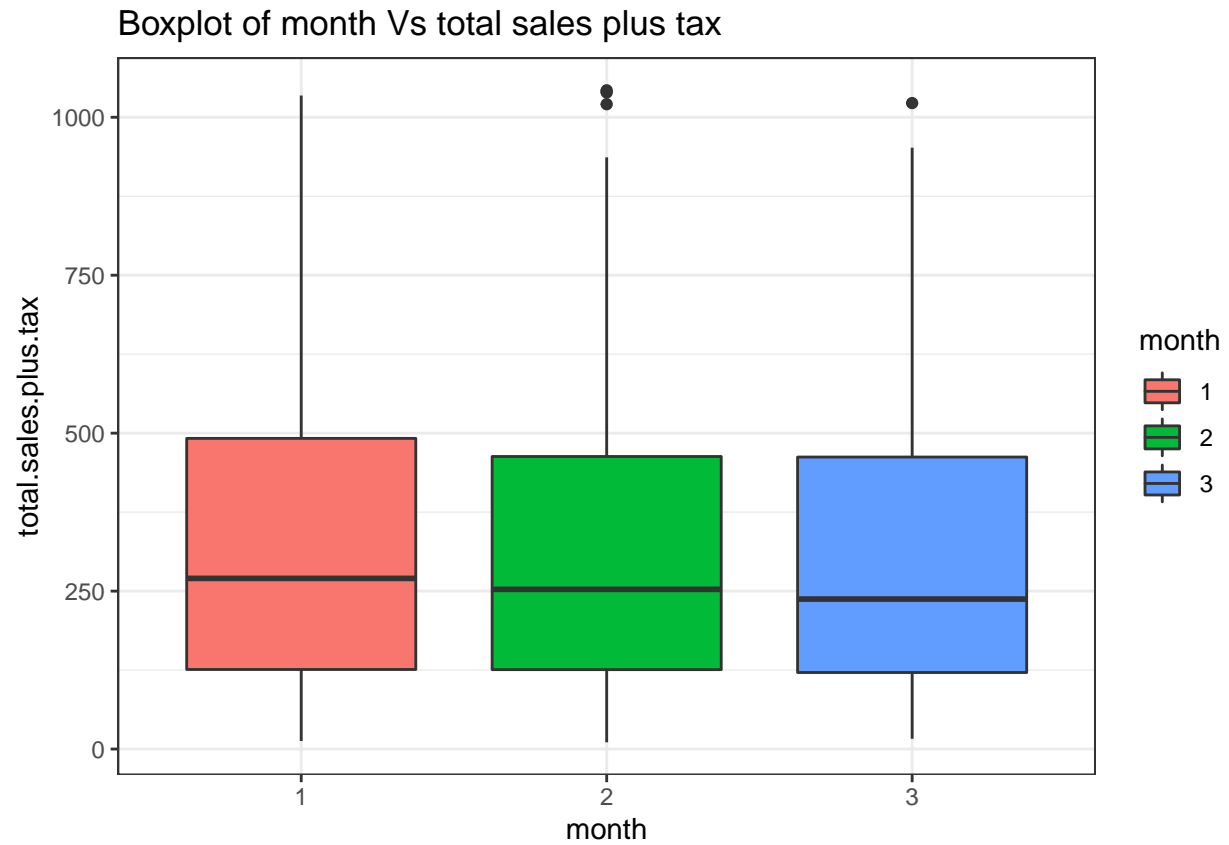
Most number of items bought is 10 followed by 9, 8, 7 all the way to 1.

```
ggplot(sales, aes(x=payment, y=total.sales.plus.tax, fill = payment)) +
  theme_bw() +
  geom_boxplot() +
  labs(title = "Boxplot of product line Vs total sales plus tax") +
  theme(axis.text.x = element_text(angle = 45))
```




On average, cash is the main mode of payment.

```
ggplot(sales, aes(x=month, y=total.sales.plus.tax, fill = month)) +  
  theme_bw() +  
  geom_boxplot() +  
  labs(title = "Boxplot of month Vs total sales plus tax") +  
  theme(axis.text.x = element_text(angle = 0))
```



On average, January brought in more total sales plus taxes.


```
ggplot(sales, aes(x=day, y=total.sales.plus.tax, fill = day)) +  
  theme_bw() +  
  geom_boxplot() +  
  labs(title = "Boxplot of day Vs total sales plus tax") +  
  theme(axis.text.x = element_text(angle = 90))
```



5.2.2 Scatter plots to determine the relationship between target variable and numerical variables


```
ggplot(sales, aes(x = unit.price, y = total.sales.plus.tax)) +
  geom_point(aes(color = "green")) +
  theme_bw() +
  labs(title = "Scatterplot of unit price vs total sales plus tax")
```

Scatterplot of unit price vs total sales plus tax



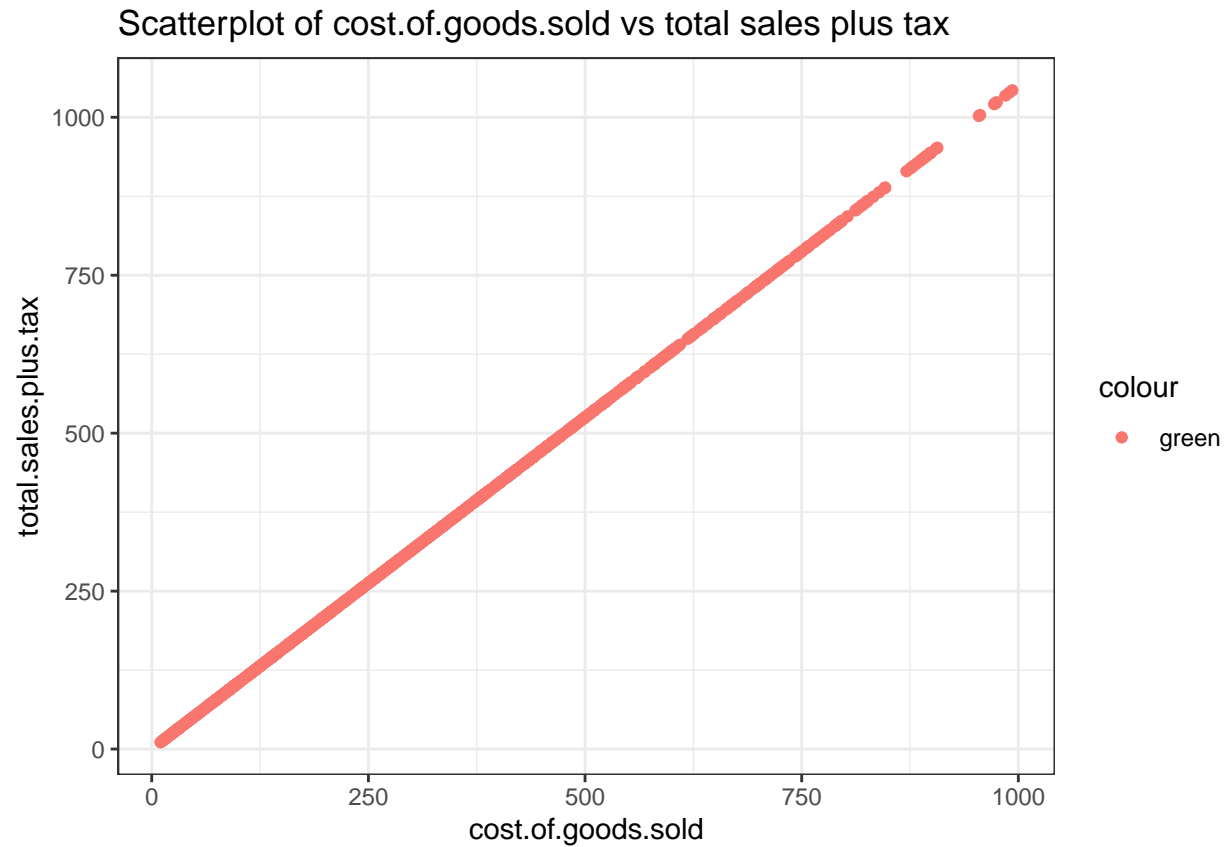
```
ggplot(sales, aes(x = tax, y = total.sales.plus.tax)) +  
  geom_point() +  
  theme_bw() +  
  labs(title = "Scatterplot of tax vs total sales plus tax")
```

Scatterplot of tax vs total sales plus tax



Tax vs sales shows a very strong positive relationship.


```
ggplot(sales, aes(x = cost.of.goods.sold, y = total.sales.plus.tax)) +  
  geom_point(aes(color = "green")) +  
  theme_bw() +  
  labs(title = "Scatterplot of cost.of.goods.sold vs total sales plus tax")
```



Another very strong relationship between cost of goods sold and total sales.

```
ggplot(sales, aes(x = gross.profit, y = total.sales.plus.tax)) +  
  geom_point(aes(color = "green")) +  
  theme_bw() +  
  labs(title = "Scatterplot of gross.profit vs total sales plus tax")
```

Scatterplot of gross.profit vs total sales plus tax



Same with the two plots above. Gross profit vs total sales shows that there is a strong positive relationship.

5.3 Data preprocessing

```
# encoding categorical variables
sales$branch <- factor(sales$branch, levels = c("A", "B", "C"), labels = c(1,2,3))
sales$gender <- factor(sales$gender, levels = c("Female", "Male"), labels = c(1,0))
sales$customer.type <- factor(sales$customer.type, levels = c("Member", "Normal"), labels = c(1,2))
sales$product.line <- factor(sales$product.line, levels = c("Health and beauty", "Electronic accessories",
"Food and beverages", "Fashion accessories"), labels = c(1,2,3,4,5,6))
sales$quantity <- as.factor(sales$quantity)
sales$payment <- factor(sales$payment, levels = c("Ewallet", "Cash", "Credit card"), labels = c(1,2,3))
glimpse(sales)
```

```
## Rows: 1,000
## Columns: 17
## $ branch          <fct> 1, 3, 1, 1, 1, 3, 1, 3, 1, 2, 2, 2, 1, 1, 1...
## $ customer.type   <fct> 1, 2, 2, 1, 2, 2, 1, 2, 1, 1, 1, 1, 2, 2, 2...
## $ gender          <fct> 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1...
## $ product.line    <fct> 1, 2, 3, 1, 4, 2, 2, 3, 1, 5, 6, 2, 2, 5, 1...
## $ unit.price      <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity        <fct> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax             <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ payment         <fct> 1, 2, 3, 1, 1, 1, 1, 1, 3, 3, 1, 2, 1, 1, 2...
## $ cost.of.goods.sold <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
```



```
## $ gross.profit      <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating            <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total.sales.plus.tax <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
## $ month             <fct> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day               <fct> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ hour              <fct> 13, 10, 13, 20, 10, 18, 14, 11, 17, 13, 18,...
## $ minute            <fct> 8, 29, 23, 33, 37, 30, 36, 38, 15, 27, 7, 3...
```

Great! All the categorical variables have been encoded.

6. Implementing the Solution

6.1 Principal Component Analysis

```
glimpse(sales)
```

```
## Rows: 1,000
## Columns: 17
## $ branch           <fct> 1, 3, 1, 1, 1, 3, 1, 3, 1, 2, 2, 2, 1, 1, 1...
## $ customer.type    <fct> 1, 2, 2, 1, 2, 2, 1, 2, 1, 1, 1, 1, 2, 2, 2...
## $ gender           <fct> 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1...
## $ product.line     <fct> 1, 2, 3, 1, 4, 2, 2, 3, 1, 5, 6, 2, 2, 5, 1...
## $ unit.price       <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity         <fct> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax              <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ payment          <fct> 1, 2, 3, 1, 1, 1, 1, 1, 3, 3, 1, 2, 1, 1, 2...
## $ cost.of.goods.sold <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.profit     <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating           <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ total.sales.plus.tax <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634....
## $ month            <fct> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day              <fct> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ hour             <fct> 13, 10, 13, 20, 10, 18, 14, 11, 17, 13, 18,...
## $ minute           <fct> 8, 29, 23, 33, 37, 30, 36, 38, 15, 27, 7, 3...
```

*# since it's unsupervised learning, we will drop the target variable(total.sales.plus.tax) before apply
first we create a copy of our cleaned dataset*

```
sales.copy <- sales
sales.df <- sales.copy[,c(1:12, 14:17)] # selecting all columns minus the target variable
glimpse(sales.df)
```

```
## Rows: 1,000
## Columns: 16
## $ branch           <fct> 1, 3, 1, 1, 1, 3, 1, 3, 1, 2, 2, 2, 1, 1, 1...
## $ customer.type    <fct> 1, 2, 2, 1, 2, 2, 1, 2, 1, 1, 1, 1, 2, 2, 2...
## $ gender           <fct> 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1...
## $ product.line     <fct> 1, 2, 3, 1, 4, 2, 2, 3, 1, 5, 6, 2, 2, 5, 1...
## $ unit.price       <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity         <fct> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10,...
## $ tax              <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ payment          <fct> 1, 2, 3, 1, 1, 1, 1, 1, 3, 3, 1, 2, 1, 1, 2...
## $ cost.of.goods.sold <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.profit     <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
```

```
## $ rating          <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ month           <fct> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day             <fct> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ hour            <fct> 13, 10, 13, 20, 10, 18, 14, 11, 17, 13, 18,...
## $ minute          <fct> 8, 29, 23, 33, 37, 30, 36, 38, 15, 27, 7, 3...
```

```
# converting the factors to numeric
```

```
sales.df$branch <- as.numeric(sales.df$branch)
sales.df$customer.type <- as.numeric(sales.df$customer.type)
sales.df$gender <- as.numeric(sales.df$gender)
sales.df$product.line <- as.numeric(sales.df$product.line)
sales.df$quantity <- as.numeric(sales.df$quantity)
sales.df$payment <- as.numeric(sales.df$payment)
sales.df$month <- as.numeric(sales.df$month)
sales.df$day <- as.numeric(sales.df$day)
sales.df$hour <- as.numeric(sales.df$hour)
sales.df$minute <- as.numeric(sales.df$minute)
glimpse(sales.df)
```


```
## Rows: 1,000
## Columns: 16
## $ branch          <dbl> 1, 3, 1, 1, 1, 3, 1, 3, 1, 2, 2, 2, 1, 1, 1...
## $ customer.type   <dbl> 1, 2, 2, 1, 2, 2, 1, 2, 1, 1, 1, 1, 2, 2, 2...
## $ gender          <dbl> 1, 1, 2, 2, 2, 2, 1, 1, 1, 1, 1, 2, 1, 2, 1...
## $ product.line    <dbl> 1, 2, 3, 1, 4, 2, 2, 3, 1, 5, 6, 2, 2, 5, 1...
## $ unit.price       <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 6...
## $ quantity        <dbl> 8, 6, 8, 9, 8, 8, 7, 2, 3, 4, 5, 5, 6, 2, 2...
## $ tax             <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ payment          <dbl> 1, 2, 3, 1, 1, 1, 1, 1, 3, 3, 1, 2, 1, 1, 2...
## $ cost.of.goods.sold <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597....
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7...
## $ gross.profit     <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085,...
## $ rating          <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2...
## $ month           <dbl> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3...
## $ day             <dbl> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 1...
## $ hour            <dbl> 4, 1, 4, 11, 1, 9, 5, 2, 8, 4, 9, 8, 1, 7, ...
## $ minute          <dbl> 9, 30, 24, 34, 38, 31, 37, 39, 16, 28, 8, 4...
```

```
# let's make a copy of the cleaned df before deleting gross.margin.percentage column that we won't need
```

```
sales.cleaned <- sales.df
```

```
sales.df$gross.margin.percentage <- NULL
```

```
corr <- cor(sales.df, method = "pearson")
corrplot(corr, method = "number", type = "lower", order = "hclust")
```



```
# applying PCA to sales.df
salesdf.pca <- prcomp(sales.df, center = TRUE, scale. = TRUE)
salesdf.pca
```

```
## Standard deviations (1, ..., p=15):
## [1] 1.902352e+00 1.094060e+00 1.083265e+00 1.051402e+00 1.028494e+00
## [6] 1.002271e+00 9.964909e-01 9.947646e-01 9.759575e-01 9.362636e-01
## [11] 9.200540e-01 8.918542e-01 6.239360e-01 4.267481e-16 1.755044e-16
##
## Rotation (n x k) = (15 x 15):
##
##          PC1          PC2          PC3          PC4
## branch      0.0273694066  0.340928572 -0.120950985  0.128779137
## customer.type -0.0149226201 -0.339573569 -0.135179876 -0.275012233
## gender      -0.0330190817 -0.520647166 -0.233365004  0.255126077
## product.line -0.0097333598  0.406034897 -0.175218710  0.300233001
## unit.price   0.3779711874  0.010913138 -0.235974297  0.168691387
## quantity    0.2143920663 -0.119904309  0.334139506 -0.310431607
## tax         0.5189385679 -0.004432186  0.008128015 -0.002731329
## payment     0.0015364151  0.397636738  0.146012221  0.140745029
## cost.of.goods.sold 0.5189385679 -0.004432186  0.008128015 -0.002731329
## gross.profit 0.5189385679 -0.004432186  0.008128015 -0.002731329
## rating      -0.0217865729  0.138902504 -0.163323789 -0.379826387
## month       -0.0159150267 -0.275084821  0.527648990  0.187287922
## day         0.0045260022 -0.162624631 -0.585795727 -0.024102676
## hour        0.0002467417 -0.195292029 -0.029978840  0.596116668
## minute     -0.0193518844  0.017272265 -0.209623749 -0.269477372
```

```

##          PC5          PC6          PC7          PC8
## branch      -0.397236394 -0.017279848 -0.42684622  0.286116629
## customer.type  0.006899145  0.580065776 -0.04654431 -0.054552999
## gender      -0.064051105 -0.272562913  0.14097483 -0.295260950
## product.line -0.157106801  0.538227869  0.18433526 -0.192602771
## unit.price   -0.039816314 -0.002339985  0.30282385 -0.008221295
## quantity     0.018999580  0.029137798 -0.53135481 -0.084137073
## tax          -0.003919154 -0.004556045  0.01486185 -0.007690004
## payment       0.366804053 -0.386908742 -0.01204285 -0.300044754
## cost.of.goods.sold -0.003919154 -0.004556045  0.01486185 -0.007690004
## gross.profit  -0.003919154 -0.004556045  0.01486185 -0.007690004
## rating       -0.273090681 -0.056250114 -0.06600763 -0.761636698
## month        -0.297676208  0.003974980  0.23180156 -0.007338774
## day          0.308497395 -0.113259870 -0.23821663  0.151258614
## hour         -0.213942348  0.017942417 -0.50162094 -0.217054741
## minute       -0.610492692 -0.363731471  0.14309078  0.208324847
##          PC9          PC10         PC11         PC12
## branch      0.53796619  0.28031075 -0.06239925  0.2389807211
## customer.type  0.52198787 -0.38706594  0.13930244 -0.0842397197
## gender      0.15689124  0.21700636  0.37036958  0.4618390344
## product.line -0.30807353 -0.18077969  0.08416226  0.4396118976
## unit.price   0.10044691  0.04868469 -0.13012445 -0.1865000405
## quantity    -0.28159550 -0.07932242  0.14726300  0.3874975218
## tax          0.01119192 -0.01034829  0.01567097 -0.0005151875
## payment      0.37819619 -0.50802066  0.09694308  0.1142223168
## cost.of.goods.sold 0.01119192 -0.01034829  0.01567097 -0.0005151875
## gross.profit  0.01119192 -0.01034829  0.01567097 -0.0005151875
## rating       0.03386215  0.18970845 -0.30834589 -0.1116389104
## month        0.13482309 -0.18831694 -0.59152925  0.2535667449
## day         -0.13009186 -0.22483767 -0.54820278  0.2658235267
## hour         -0.15140907 -0.23102958  0.02263424 -0.4285155469
## minute       -0.16318895 -0.49938469  0.19408286 -0.0143798523
##          PC13         PC14         PC15
## branch      0.0257384233 -2.382931e-17  4.525329e-18
## customer.type  0.0101291399 -6.448407e-17 -1.255668e-16
## gender      -0.0258103131 -8.810457e-17  2.539151e-17
## product.line -0.0286333690 -3.103373e-17 -1.155330e-17
## unit.price   0.7844379686 -2.997054e-16  1.041429e-16
## quantity     0.4284696429 -2.615067e-16  9.632691e-17
## tax          -0.2512938678 -8.150962e-01  4.780050e-02
## payment      0.0526916416 -1.167458e-16  1.062647e-16
## cost.of.goods.sold -0.2512938678  4.489445e-01  6.819937e-01
## gross.profit  -0.2512938678  3.661516e-01 -7.297942e-01
## rating       -0.0492164258  3.800031e-17  3.665228e-17
## month        -0.0181274544  6.084707e-17  9.785773e-17
## day          -0.0592122830  1.307555e-16  1.233559e-16
## hour         -0.0184473676  4.306750e-17 -2.027034e-17
## minute       -0.0004861354 -1.445673e-17  2.158154e-17

```

```
names(salesdf.pca)
```


```
## [1] "sdev"      "rotation" "center"   "scale"    "x"
```

```
summary(salesdf.pca)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  1.9024 1.0941 1.08327 1.0514 1.02849 1.00227 0.9965
## Proportion of Variance 0.2413 0.0798 0.07823 0.0737 0.07052 0.06697 0.0662
## Cumulative Proportion 0.2413 0.3211 0.39929 0.4730 0.54351 0.61048 0.6767
##          PC8      PC9      PC10      PC11      PC12      PC13      PC14
## Standard deviation  0.99476 0.9760 0.93626 0.92005 0.89185 0.62394 4.267e-16
## Proportion of Variance 0.06597 0.0635 0.05844 0.05643 0.05303 0.02595 0.000e+00
## Cumulative Proportion 0.74265 0.8062 0.86459 0.92102 0.97405 1.00000 1.000e+00
##          PC15
## Standard deviation  1.755e-16
## Proportion of Variance 0.000e+00
## Cumulative Proportion 1.000e+00
```

```
# visualization of the PCA
```

```
ggbiplot(salesdf.pca, obs.scale = 1, var.scale = 0.5, circle = TRUE, ellipse = TRUE, varname.size = 5) +
  theme_bw() +
  theme(legend.direction = 'horizontal', legend.position = 'top') +
  geom_point(size = 2, color = 'light blue')
```



From the plot, the significant variables affecting our target variable is cost of goods sold, followed by gross profit, the branch, quantity and mode of payment. If we were to create a model we would pick these variables as our top choices because they influence the target variable the most.

6.2 Feature Selection

6.2.1 Filter Methods

```
# with filter method, we use correlation matrix to filter out the most correlated variables. If two var  
# let's create a copy of the dataset to use for feature selection
```

```
sales.fs <- sales.cleaned  
# step 1: calculating the correlation matrix for the dataset  
sales.corr <- cor(sales.fs)
```

```
## Warning in cor(sales.fs): the standard deviation is zero
```

```
sales.corr
```

```
##           branch customer.type      gender product.line  
## branch      1.00000000 -0.019607869 -0.056317558  0.03290205  
## customer.type -0.01960787  1.000000000  0.039996160 -0.02510945  
## gender      -0.05631756  0.039996160  1.000000000 -0.06612647  
## product.line  0.03290205 -0.025109450 -0.066126475  1.00000000  
## unit.price    0.02820244 -0.020237875  0.015444630  0.03842765  
## quantity    -0.02001652  0.004744509 -0.061971894 -0.06977620  
## tax          0.04104666 -0.019670283 -0.049450989 -0.01854396  
## payment      0.02055290 -0.069286242 -0.049514182  0.01051098  
## cost.of.goods.sold 0.04104666 -0.019670283 -0.049450989 -0.01854396  
## gross.margin.percentage NA          NA          NA          NA  
## gross.profit   0.04104666 -0.019670283 -0.049450989 -0.01854396  
## rating        0.01023848  0.018888672  0.004800208  0.02339096  
## month        -0.03530092  0.005972443  0.027533609 -0.04701346  
## day          -0.01308653  0.034124208  0.051156850 -0.02332870  
## hour         0.03300711 -0.018893298  0.084081139  0.03691312  
## minute       0.03837833 -0.012909043  0.009257593 -0.01014963  
##           unit.price  quantity      tax      payment  
## branch      0.028202440 -0.020016524  0.041046665  0.020552896  
## customer.type -0.020237875  0.004744509 -0.019670283 -0.069286242  
## gender      0.015444630 -0.061971894 -0.049450989 -0.049514182  
## product.line 0.038427649 -0.069776204 -0.018543956  0.010510982  
## unit.price   1.000000000  0.008127624  0.633962089 -0.019637884  
## quantity     0.008127624  1.000000000  0.357573247 -0.029577901  
## tax          0.633962089  0.357573247  1.000000000  0.008823723  
## payment     -0.019637884 -0.029577901  0.008823723  1.000000000  
## cost.of.goods.sold 0.633962089  0.357573247  1.000000000  0.008823723  
## gross.margin.percentage NA          NA          NA          NA  
## gross.profit  0.633962089  0.357573247  1.000000000  0.008823723  
## rating      -0.008777507  0.017240731 -0.036441705  0.013001094  
## month      -0.027387186  0.020515373 -0.022301340 -0.022555784  
## day        0.057020896 -0.024342312 -0.002514770 -0.028627647  
## hour       0.008242210  0.015210429 -0.002770440 -0.013989800  
## minute    -0.006868818 -0.024797102 -0.027479899 -0.050585696  
##           cost.of.goods.sold gross.margin.percentage gross.profit  
## branch      0.041046665          NA  0.041046665  
## customer.type -0.019670283          NA -0.019670283  
## gender      -0.049450989          NA -0.049450989  
## product.line -0.018543956          NA -0.018543956  
## unit.price   0.633962089          NA  0.633962089  
## quantity     0.357573247          NA  0.357573247  
## tax          1.000000000          NA  1.000000000
```

```
## payment 0.008823723 NA 0.008823723
## cost.of.goods.sold 1.000000000 NA 1.000000000
## gross.margin.percentage NA 1 NA
## gross.profit 1.000000000 NA 1.000000000
## rating -0.036441705 NA -0.036441705
## month -0.022301340 NA -0.022301340
## day -0.002514770 NA -0.002514770
## hour -0.002770440 NA -0.002770440
## minute -0.027479899 NA -0.027479899
## rating month day hour
## branch 0.010238476 -0.035300925 -0.013086533 0.03300711
## customer.type 0.018888672 0.005972443 0.034124208 -0.01889330
## gender 0.004800208 0.027533609 0.051156850 0.08408114
## product.line 0.023390962 -0.047013462 -0.023328697 0.03691312
## unit.price -0.008777507 -0.027387186 0.057020896 0.00824221
## quantity 0.017240731 0.020515373 -0.024342312 0.01521043
## tax -0.036441705 -0.022301340 -0.002514770 -0.00277044
## payment 0.013001094 -0.022555784 -0.028627647 -0.01398980
## cost.of.goods.sold -0.036441705 -0.022301340 -0.002514770 -0.00277044
## gross.margin.percentage NA NA NA NA
## gross.profit -0.036441705 -0.022301340 -0.002514770 -0.00277044
## rating 1.000000000 -0.042880374 -0.007075821 -0.03058764
## month -0.042880374 1.000000000 -0.118996386 0.04376174
## day -0.007075821 -0.118996386 1.000000000 0.02066810
## hour -0.030587644 0.043761744 0.020668100 1.00000000
## minute 0.050558480 -0.006553809 0.012645496 -0.02538363
## minute
## branch 0.038378328
## customer.type -0.012909043
## gender 0.009257593
## product.line -0.010149626
## unit.price -0.006868818
## quantity -0.024797102
## tax -0.027479899
## payment -0.050585696
## cost.of.goods.sold -0.027479899
## gross.margin.percentage NA
## gross.profit -0.027479899
## rating 0.050558480
## month -0.006553809
## day 0.012645496
## hour -0.025383629
## minute 1.000000000
```

We see NA for gross margin percentage. let's delete the column and redo the correlation for the remaining columns.

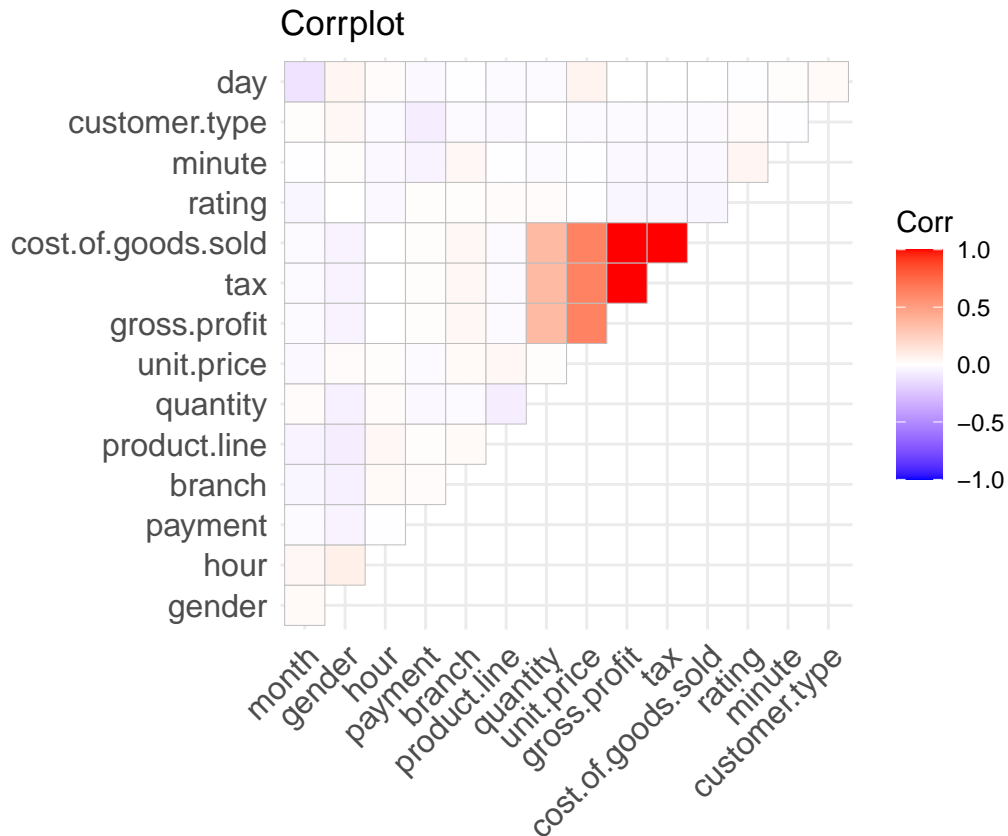
```
sales.fs$gross.margin.percentage <- NULL
salesfs.corr <- cor(sales.fs)
salesfs.corr
```

```
## branch customer.type gender product.line
## branch 1.00000000 -0.019607869 -0.056317558 0.03290205
## customer.type -0.01960787 1.000000000 0.039996160 -0.02510945
## gender -0.05631756 0.039996160 1.000000000 -0.06612647
```

## product.line	0.03290205	-0.025109450	-0.066126475	1.00000000
## unit.price	0.02820244	-0.020237875	0.015444630	0.03842765
## quantity	-0.02001652	0.004744509	-0.061971894	-0.06977620
## tax	0.04104666	-0.019670283	-0.049450989	-0.01854396
## payment	0.02055290	-0.069286242	-0.049514182	0.01051098
## cost.of.goods.sold	0.04104666	-0.019670283	-0.049450989	-0.01854396
## gross.profit	0.04104666	-0.019670283	-0.049450989	-0.01854396
## rating	0.01023848	0.018888672	0.004800208	0.02339096
## month	-0.03530092	0.005972443	0.027533609	-0.04701346
## day	-0.01308653	0.034124208	0.051156850	-0.02332870
## hour	0.03300711	-0.018893298	0.084081139	0.03691312
## minute	0.03837833	-0.012909043	0.009257593	-0.01014963
##	unit.price	quantity	tax	payment
## branch	0.028202440	-0.020016524	0.041046665	0.020552896
## customer.type	-0.020237875	0.004744509	-0.019670283	-0.069286242
## gender	0.015444630	-0.061971894	-0.049450989	-0.049514182
## product.line	0.038427649	-0.069776204	-0.018543956	0.010510982
## unit.price	1.000000000	0.008127624	0.633962089	-0.019637884
## quantity	0.008127624	1.000000000	0.357573247	-0.029577901
## tax	0.633962089	0.357573247	1.000000000	0.008823723
## payment	-0.019637884	-0.029577901	0.008823723	1.000000000
## cost.of.goods.sold	0.633962089	0.357573247	1.000000000	0.008823723
## gross.profit	0.633962089	0.357573247	1.000000000	0.008823723
## rating	-0.008777507	0.017240731	-0.036441705	0.013001094
## month	-0.027387186	0.020515373	-0.022301340	-0.022555784
## day	0.057020896	-0.024342312	-0.002514770	-0.028627647
## hour	0.008242210	0.015210429	-0.002770440	-0.013989800
## minute	-0.006868818	-0.024797102	-0.027479899	-0.050585696
##	cost.of.goods.sold	gross.profit	rating	month
## branch	0.041046665	0.041046665	0.010238476	-0.035300925
## customer.type	-0.019670283	-0.019670283	0.018888672	0.005972443
## gender	-0.049450989	-0.049450989	0.004800208	0.027533609
## product.line	-0.018543956	-0.018543956	0.023390962	-0.047013462
## unit.price	0.633962089	0.633962089	-0.008777507	-0.027387186
## quantity	0.357573247	0.357573247	0.017240731	0.020515373
## tax	1.000000000	1.000000000	-0.036441705	-0.022301340
## payment	0.008823723	0.008823723	0.013001094	-0.022555784
## cost.of.goods.sold	1.000000000	1.000000000	-0.036441705	-0.022301340
## gross.profit	1.000000000	1.000000000	-0.036441705	-0.022301340
## rating	-0.036441705	-0.036441705	1.000000000	-0.042880374
## month	-0.022301340	-0.022301340	-0.042880374	1.000000000
## day	-0.002514770	-0.002514770	-0.007075821	-0.118996386
## hour	-0.002770440	-0.002770440	-0.030587644	0.043761744
## minute	-0.027479899	-0.027479899	0.050558480	-0.006553809
##	day	hour	minute	
## branch	-0.013086533	0.03300711	0.038378328	
## customer.type	0.034124208	-0.01889330	-0.012909043	
## gender	0.051156850	0.08408114	0.009257593	
## product.line	-0.023328697	0.03691312	-0.010149626	
## unit.price	0.057020896	0.00824221	-0.006868818	
## quantity	-0.024342312	0.01521043	-0.024797102	
## tax	-0.002514770	-0.00277044	-0.027479899	
## payment	-0.028627647	-0.01398980	-0.050585696	
## cost.of.goods.sold	-0.002514770	-0.00277044	-0.027479899	


```
## gross.profit      -0.002514770 -0.00277044 -0.027479899
## rating            -0.007075821 -0.03058764  0.050558480
## month             -0.118996386  0.04376174 -0.006553809
## day               1.000000000  0.02066810  0.012645496
## hour              0.020668100  1.000000000 -0.025383629
## minute            0.012645496 -0.02538363  1.000000000
```

```
sales.fs %>% cor %>% ggcorrplot(method = "square", lab = FALSE, type = "upper", hc.order = TRUE, title
```



```
# finding highly correlated attributes
```

```
high.corr <- findCorrelation(salesfs.corr, cutoff = 0.75)
names(sales.fs[, high.corr])
```

```
## [1] "tax" "cost.of.goods.sold"
```

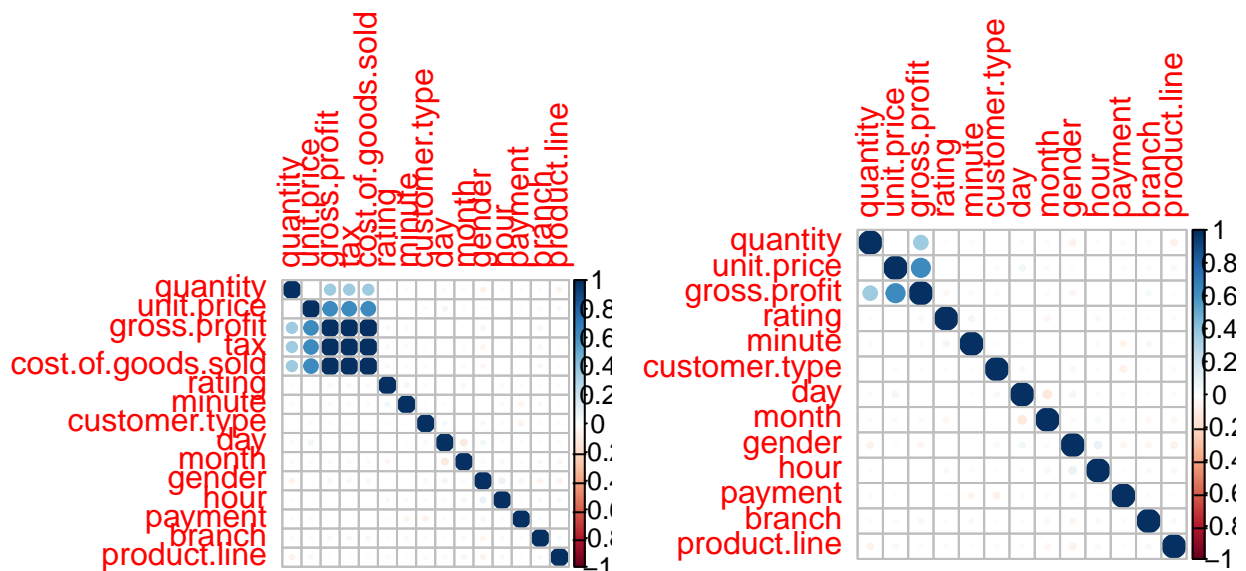
```
# removing highly correlated features to remove redundant features
```

```
sales.fs2 <- sales.fs[, -high.corr]
head(sales.fs2)
```

```
##   branch customer.type gender product.line unit.price quantity payment
## 1      1             1      1             1      74.69         8        1
## 2      3             2      1             2      15.28         6        2
## 3      1             2      2             3      46.33         8        3
## 4      1             1      2             1      58.22         9        1
## 5      1             2      2             4      86.31         8        1
## 6      3             2      2             2      85.39         8        1
```

```
##   gross.profit rating month day hour minute
## 1      26.1415   9.1     1   5    4      9
## 2       3.8200   9.6     3   8    1     30
## 3      16.2155   7.4     3   3    4     24
## 4      23.2880   8.4     1  27   11     34
## 5      30.2085   5.3     2   8    1     38
## 6      29.8865   4.1     3  25    9     31
```

```
# we can now compare the two correlations, before and after removing redundant variables
par(mfrow = c(1, 2), pty = "s")
corrplot(salesfs.corr, order = "hclust")
corrplot(cor(sales.fs2), order = "hclust")
```



6.2.2 Wrapper Methods

```
# making a copy of df to use for wrapper method
```

```
sales.wm <- sales.df
glimpse(sales.wm)
```

```
## Rows: 1,000
## Columns: 15
## $ branch      <dbl> 1, 3, 1, 1, 1, 3, 1, 3, 1, 2, 2, 2, 1, 1, 1, 2, ...
## $ customer.type <dbl> 1, 2, 2, 1, 2, 2, 1, 2, 1, 1, 1, 1, 2, 2, 2, 1, ...
## $ gender      <dbl> 1, 1, 2, 2, 2, 2, 1, 1, 1, 1, 1, 2, 1, 2, 1, 1, ...
## $ product.line <dbl> 1, 2, 3, 1, 4, 2, 2, 3, 1, 5, 6, 2, 2, 5, 1, 4, ...
## $ unit.price   <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 68.84, ...
```

```
## $ quantity      <dbl> 8, 6, 8, 9, 8, 8, 7, 2, 3, 4, 5, 5, 6, 2, 2, 7, ...
## $ tax            <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085, 29.8...
## $ payment        <dbl> 1, 2, 3, 1, 1, 1, 1, 1, 3, 3, 1, 2, 1, 1, 2, 2, ...
## $ cost.of.goods.sold <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597.73, 4...
## $ gross.profit    <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085, 29.8...
## $ rating          <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2, 5.9...
## $ month           <dbl> 1, 3, 3, 1, 2, 3, 2, 2, 1, 2, 2, 3, 2, 2, 3, 1, ...
## $ day             <dbl> 5, 8, 3, 27, 8, 25, 25, 24, 10, 20, 6, 9, 12, 7,...
## $ hour            <dbl> 4, 1, 4, 11, 1, 9, 5, 2, 8, 4, 9, 8, 1, 7, 10, 7...
## $ minute          <dbl> 9, 30, 24, 34, 38, 31, 37, 39, 16, 28, 8, 4, 26,...
```

we begin with a sequential forward greedy search

```
wrapper = clustvarsel(sales.wm, G = 1:15)
wrapper
```

```
## -----
## Variable selection for Gaussian model-based clustering
## Stepwise (forward/backward) greedy search
## -----
##
## Variable proposed Type of step BICclust Model G BICdiff Decision
##          hour      Add -4267.096      E 10  901.6667 Accepted
##          branch    Add -6237.760     VEV  4  477.3458 Accepted
##          month     Add -9395.372     EEI  5 -666.9596 Rejected
##          branch    Remove -4267.096      E 10  477.3458 Rejected
##
## Selected subset: hour, branch
```

From the search two variables were selected: hour and branch, the rest were rejected.

the variables identified to use for clustering are: hour and branch. Let's proceed to the modeling part


```
selected = sales.wm[, wrapper$subset] # choosing the selected variables
glimpse(selected)
```

```
## Rows: 1,000
## Columns: 2
## $ hour      <dbl> 4, 1, 4, 11, 1, 9, 5, 2, 8, 4, 9, 8, 1, 7, 10, 7, 2, 1, 9, 6...
## $ branch    <dbl> 1, 3, 1, 1, 1, 3, 1, 3, 1, 2, 2, 2, 1, 1, 1, 2, 1, 1, 1, 2, ...
```

```
model = Mclust(selected, G = 1:15) #building the model
summary(model)
```

```
## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
##
## Mclust VEV (ellipsoidal, equal shape) model with 4 components:
##
## log-likelihood    n df      BIC      ICL
##      -3049.802 1000 20 -6237.76 -6297.175
##
## Clustering table:
##    1  2  3  4
## 137 327 203 333
```

```
# plotting the model
plot(model,c("classification"))
```



6.2.3 Embedded methods

embedded method is know to be suitable for very high dimensional dataset. Ours has 15 columns after c
we first make a copy to use for this implementation part

```
sales.em <- sales.df
set.seed(2)
em.model <- ewkm(sales.em[1:15], 3, maxiter = 1000)
em.model$size
```

```
## [1] 379 231 390
```

```
em.model$centers
```

```
##      branch customer.type  gender product.line unit.price quantity      tax
## 1 2.010554      1.490765 1.000000   3.672823   46.72037 5.076517 10.325594
## 2 2.038961      1.467532 1.471861   3.484848   80.65697 6.623377 33.299985
## 3 1.935897      1.525641 2.000000   3.530769   49.57269 4.802564  9.676083
##  payment cost.of.goods.sold gross.profit  rating  month  day  hour
## 1 2.010554          206.5119   10.325594 7.000792 1.992084 14.90765 5.701847
## 2 1.982684          665.9997   33.299985 6.859740 1.948052 15.16883 5.718615
## 3 1.912821          193.5217    9.676083 7.012308 2.020513 15.64615 6.225641
##      minute
## 1 30.98417
```

```
## 2 29.83550
## 3 31.95641
```

```
names(em.model)
```

```
## [1] "cluster"      "centers"      "totss"        "withinss"
## [5] "tot.withinss" "betweenss"    "size"         "iterations"
## [9] "total.iterations" "restarts"     "weights"
```


```
round(em.model$weights*100,2)
```

```
##   branch customer.type gender product.line unit.price quantity tax payment
## 1      0           0.00  99.99          0          0          0  0         0
## 2      0          51.51  48.48          0          0          0  0         0
## 3      0           0.00  99.99          0          0          0  0         0
##   cost.of.goods.sold gross.profit rating month day hour minute
## 1                  0           0      0      0  0  0         0
## 2                  0           0      0      0  0  0         0
## 3                  0           0      0      0  0  0         0
```

In cluster 1 and 3 gender carry's the most weight, while in cluster 2, customer type and gender carry about the same weight.

```
# plotting cluster for 1st and second PC
```

```
clusplot(sales.em[1:15], em.model$cluster, color = TRUE, shade = TRUE, lines = 1, main = "Clustering analysis for sales data")
```



These two components explain 32.11 % of the point variability.

7. Follow up questions: did we have the right data? Yes we did.